

# STATS 507

# Data Analysis in Python

Week 10: PyTorch (Part 2)

# Recap: What is PyTorch?

It's a Python-based scientific computing package targeted at two sets of audiences:

1. A replacement for NumPy to use the power of GPUs
2. a deep learning research platform that provides maximum flexibility and speed

# Recap: Resizing


Resizing: If you want to resize/reshape tensor, you can use `torch.view`:

```
[14] x = torch.randn(4, 4)
      y = x.view(16)
      z = x.view(-1, 8) # the size -1 is inferred from other dimensions
      print(x.size(), y.size(), z.size())
```

 `torch.Size([4, 4]) torch.Size([16]) torch.Size([2, 8])`

If you have a one element tensor, use `.item()` to get the value as a Python number

```
▶ x = torch.randn(1)
  print(x)
  print(x.item())
```

 `tensor([-0.8768])`  
`-0.8767884373664856`

## Recap: GPU speedup on Google colab

```
[1] import torch  
import time
```

```
[2] x_cpu = torch.randn(120000, 10000)  
y_cpu = torch.randn(10000, 1)
```

```
[3] x_gpu = x_cpu.cuda()  
y_gpu = y_cpu.cuda()
```

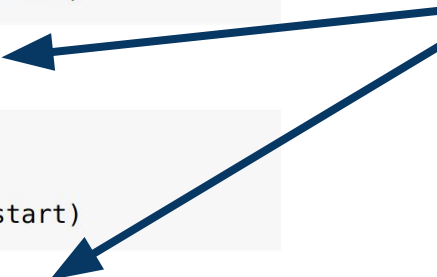
```
[4] start = time.time()  
x_cpu.mm(y_cpu)  
print(time.time() - start)
```

0.39958739280700684

```
▶ start = time.time()  
x_gpu.mm(y_gpu)  
print(time.time() - start)
```

📄 0.004973173141479492

Matrix-vector multiplication on the GPU is nearly 100x faster!



# Recap: Automatic Differentiation

```
[2] import torch
```

Create a tensor and set requires\_grad=True to track computation with it

```
[3] x = torch.ones(2, 2, requires_grad=True)
    print(x)
```

```
tensor([[1., 1.],
        [1., 1.]], requires_grad=True)
```

Do an operation of tensor:

```
[4] y = x + 2
    print(y)
```

```
tensor([[3., 3.],
        [3., 3.]], grad_fn=<AddBackward0>)
```

y was created as a result of an operation, so it has a grad\_fn.

```
[5] print(y.grad_fn)
```

```
<AddBackward0 object at 0x7f4f72142898>
```

Do more operations on y

```
[6] z = y * y * 3
    out = z.mean()
```

```
print(z, out)
```

```
tensor([[27., 27.],
        [27., 27.]], grad_fn=<MulBackward0>) tensor(27.,
```

# Recap: Automatic Differentiation

Let's backprop now Because `out` contains a single scalar, `out.backward()` is equivalent to `out.backward(torch.tensor(1))`.

```
[8] out.backward()
```

```
print gradients d(out)/dx
```

```
[9] print(x.grad)
```

```
tensor([[4.5000, 4.5000],  
        [4.5000, 4.5000]])
```

You should have got a matrix of 4.5. Let's call the `out Tensor "o"`. We have that  $o = \frac{1}{4} \sum_i z_i$ ,  $z_i = 3(x_i + 2)^2$  and  $z_i|_{x_i=1} = 27$ . Therefore,  $\frac{\partial o}{\partial x_i} = \frac{3}{2}(x_i + 2)$ , hence  $\frac{\partial o}{\partial x_i}|_{x_i=1} = \frac{9}{2} = 4.5$ .

You can do many crazy things with autograd!

# Gradient Descent & Stochastic Gradient Descent (SGD)

# Numerical Optimization

Suppose we want to minimize


$$f(x) = x^2 + x + 1$$

Easy!

$$\nabla f(x) = 0$$

$$\implies 2x + 1 = 0$$

$$\implies x = -1/2$$





# Numerical Optimization

Now consider


$$f(x) = -x^2 + 10x \log(x)$$

We know

$$\nabla f(x) = -2x + 10 \log(x) + 10$$

But what value of  $x$  sets

$$\nabla f(x) = 0 \quad ?$$




# Iterative gradient-based solvers

In [82]:

```
1 import torch
2
3 def f(x):
4     return -x**2 + 10 * x * x.log()
5
6 x = torch.ones(1, requires_grad=True)
7 y = f(x)
8 y.backward()
9 x.grad
```


Out[82]: tensor([8.])



# Gradient Descent

```
1 for i in range(20):
2     x.grad.zero_()
3     y = f(x)
4     y.backward()
5     print("x=%.3f  fx=%.3f  dfdx=%.3f" % (x, y, x.grad))
6     step_size = 0.02
7     with torch.no_grad():
8         x -= step_size * x.grad
```

x=1.000	fx=-1.000	dfdx=8.000
x=0.840	fx=-2.170	dfdx=6.576
x=0.708	fx=-2.944	dfdx=5.137
x=0.606	fx=-3.404	dfdx=3.775
x=0.530	fx=-3.645	dfdx=2.595
x=0.478	fx=-3.756	dfdx=1.669
x=0.445	fx=-3.801	dfdx=1.012
x=0.425	fx=-3.817	dfdx=0.587
x=0.413	fx=-3.823	dfdx=0.330
x=0.406	fx=-3.824	dfdx=0.182
x=0.403	fx=-3.825	dfdx=0.099
x=0.401	fx=-3.825	dfdx=0.054
x=0.400	fx=-3.825	dfdx=0.029
x=0.399	fx=-3.825	dfdx=0.016
x=0.399	fx=-3.825	dfdx=0.008
x=0.399	fx=-3.825	dfdx=0.005
x=0.398	fx=-3.825	dfdx=0.002
x=0.398	fx=-3.825	dfdx=0.001
x=0.398	fx=-3.825	dfdx=0.001
x=0.398	fx=-3.825	dfdx=0.000




### Practice Problem 3:

Minimizing our function with gradient descent

# Stochastic Gradient Descent

```
1 def noisy_f(x):
2     return -x**2 + 10 * x * x.log() + 0.1 * x * torch.randn(1)
3
4 x = torch.ones(1, requires_grad=True)
5 for i in range(1, 20):
6     y = noisy_f(x)
7     y.backward()
8     print("x=%0.3f  fx=%0.3f  dfdx=%0.3f" % (x, f(x), x.grad))
9     step_size = 0.1 / (i**1.5)
10    with torch.no_grad():
11        x -= step_size * x.grad
12    x.grad.zero_()
13
```


```
x=1.000  fx=-1.000  dfdx=7.844
x=0.216  fx=-3.355  dfdx=-5.812
x=0.421  fx=-3.819  dfdx=0.553
x=0.410  fx=-3.824  dfdx=0.268
x=0.407  fx=-3.824  dfdx=0.167
x=0.406  fx=-3.825  dfdx=0.053
x=0.405  fx=-3.825  dfdx=0.103
x=0.405  fx=-3.825  dfdx=0.141
x=0.404  fx=-3.825  dfdx=0.212
x=0.403  fx=-3.825  dfdx=-0.020
x=0.403  fx=-3.825  dfdx=0.285
x=0.403  fx=-3.825  dfdx=0.233
x=0.402  fx=-3.825  dfdx=0.021
x=0.402  fx=-3.825  dfdx=0.116
x=0.402  fx=-3.825  dfdx=0.107
x=0.402  fx=-3.825  dfdx=0.132
x=0.401  fx=-3.825  dfdx=0.070
x=0.401  fx=-3.825  dfdx=0.222
x=0.401  fx=-3.825  dfdx=0.120
```



# PyTorch Optimizers

```
1 x = torch.ones(1, requires_grad=True)
2
3 optimizer = torch.optim.SGD([x], lr=0.02)
4
5 for i in range(1, 15):
6     y = noisy_f(x)
7     y.backward()
8     print("x=%.3f  fx=%.3f  dfdx=%.3f" % (x, f(x), x.grad))
9     optimizer.step()
10    x.grad.zero_()
11
```

x=1.000	fx=-1.000	dfdx=8.171
x=0.837	fx=-2.193	dfdx=6.391
x=0.709	fx=-2.942	dfdx=5.235
x=0.604	fx=-3.410	dfdx=3.756
x=0.529	fx=-3.648	dfdx=2.561
x=0.478	fx=-3.757	dfdx=1.597
x=0.446	fx=-3.800	dfdx=0.977
x=0.426	fx=-3.816	dfdx=0.587
x=0.414	fx=-3.822	dfdx=0.409
x=0.406	fx=-3.824	dfdx=0.107
x=0.404	fx=-3.825	dfdx=0.081
x=0.403	fx=-3.825	dfdx=0.218
x=0.398	fx=-3.825	dfdx=-0.066
x=0.400	fx=-3.825	dfdx=0.087



# Loss functions as objective functions

```
1  n = 10000
2  p = 5
3  dataset_x = torch.randn(n, p) # a synthetic dataset
4
5  beta = torch.zeros(p, 1)
6  beta[2] = -1.0
7  beta[4] = 3.0
8
9  dataset_y = torch.mm(dataset_x, beta) # and a response
10
11 def loss(x, y, beta_hat):
12     yhat = torch.mm(x, beta_hat)
13     return ((yhat - y)**2).mean()
14
15 beta_hat = torch.ones((p, 1), requires_grad=True)
16
17 optimizer = torch.optim.SGD([beta_hat,], lr=2e-1)
18
19 for i in range(1, 15):
20     mse = loss(dataset_x, dataset_y, beta_hat)
21     mse.backward()
22     print("mse = %.3f" % (mse.item()))
23     optimizer.step()
24     beta_hat.grad.zero_()
```

# Loss functions as objective functions

```
1 n = 10000
2 p = 5
3 dataset_x = torch.randn(n, p) # a synthetic dataset
4
5 beta = torch.zeros(p, 1)
6 beta[2] = -1.0
7 beta[4] = 3.0
8
9 dataset_y = torch.mm(dataset_x, beta) # and a response
10
11 def loss(x, y, beta_hat):
12     yhat = torch.mm(x, beta_hat)
13     return ((yhat - y)**2).mean()
14
15 beta_hat = torch.ones((p, 1), requires_grad=True)
16
17 optimizer = torch.optim.SGD([beta_hat,], lr=2e-1)
18
19 for i in range(1, 15):
20     mse = loss(dataset_x, dataset_y, beta_hat)
21     mse.backward()
22     print("mse = %.3f" % (mse.item()))
23     optimizer.step()
24     beta_hat.grad.zero_()
```

```
mse = 11.019
mse = 3.956
mse = 1.421
mse = 0.511
mse = 0.184
mse = 0.066
mse = 0.024
mse = 0.009
mse = 0.003
mse = 0.001
mse = 0.000
mse = 0.000
mse = 0.000
mse = 0.000
```

1	beta_hat
tensor([[ 3.8991e-04], [ 6.4804e-04], [-9.9832e-01], [ 9.9378e-04], [ 2.9984e+00]], requires_grad=True)	



# Stochastic optimization with minibatches

```
1 n = 10000
2 p = 5
3 dataset_x = torch.randn(n, p) # a synthetic dataset
4
5 beta = torch.zeros(p, 1)
6 beta[2] = -1.0
7 beta[4] = 3.0
8
9 dataset_y = torch.mm(dataset_x, beta) # and a response
10
11 def loss(x, y, beta_hat):
12     yhat = torch.mm(x, beta_hat)
13     return ((yhat - y)**2).mean()
14
15 beta_hat = torch.ones((p, 1), requires_grad=True)
16
17 optimizer = torch.optim.SGD([beta_hat,], lr=2e-1)
18
19 for i in range(1, 15):
20     minibatch = torch.randint(n, (128,))
21     mse = loss(dataset_x[minibatch], dataset_y[minibatch], beta_hat)
22     mse.backward()
23     print("mse = %.3f" % (mse.item()))
24     optimizer.step()
25     beta_hat.grad.zero_()
```

```
mse = 12.523
mse = 3.308
mse = 1.214
mse = 0.388
mse = 0.186
mse = 0.055
mse = 0.030
mse = 0.007
mse = 0.002
mse = 0.001
mse = 0.000
mse = 0.000
mse = 0.000
mse = 0.000
```

```
: 1 beta_hat
: tensor([[ 2.4282e-04],
          [ 7.5062e-04],
          [-9.9902e-01],
          [ 1.2907e-03],
          [ 2.9988e+00]], requires_grad=True)
```

# Datasets and DataLoaders

```
1 from torch.utils.data import TensorDataset, DataLoader
2
3 tt_split = int(n * .9)
4 train_ds = TensorDataset(dataset_x[:tt_split], dataset_y[:tt_split])
5 test_ds = TensorDataset(dataset_x[tt_split:], dataset_y[tt_split:])
6
7 train_dl = DataLoader(train_ds, batch_size=128, shuffle=True, num_workers=0)
8 test_dl = DataLoader(test_ds, batch_size=128, shuffle=False, num_workers=0)
9
10 beta_hat = torch.ones((p, 1), requires_grad=True)
11 optimizer = torch.optim.SGD([beta_hat,], lr=2e-1)
12
13 for i, (x, y) in enumerate(train_dl):
14     mse = loss(x, y, beta_hat)
15     mse.backward()
16     print("mse = %.3f" % (mse.item()))
17     optimizer.step()
18     beta_hat.grad.zero_()
```

```
mse = 9.103
mse = 4.647
mse = 1.902
mse = 0.829
mse = 0.231
mse = 0.104
mse = 0.030
```

# MNIST dataset

```
1 import torchvision
2 import torchvision.transforms as transforms
3
4 train_ds = torchvision.datasets.MNIST(root='.',
5                                     train=True,
6                                     transform=transforms.ToTensor(),
7                                     download=True)
8
9 test_ds = torchvision.datasets.MNIST(root='.',
10                                    train=False,
11                                    transform=transforms.ToTensor())
```

0.1%

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to ./MNIST/raw/train-images-idx3-ubyte.gz

100.1%

Extracting ./MNIST/raw/train-images-idx3-ubyte.gz to ./MNIST/raw

113.5%

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to ./MNIST/raw/train-labels-idx1-ubyte.gz

Extracting ./MNIST/raw/train-labels-idx1-ubyte.gz to ./MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to ./MNIST/raw/t10k-images-idx3-ubyte.gz

180.4%

Extracting ./MNIST/raw/t10k-images-idx3-ubyte.gz to ./MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to ./MNIST/raw/t10k-labels-idx1-ubyte.gz

Extracting ./MNIST/raw/t10k-labels-idx1-ubyte.gz to ./MNIST/raw

Processing...

Done!

# Exploring MNIST

```
1 x, y = train_ds[0]
```

```
1 y
```

```
5
```

```
1 x.shape
```


```
torch.Size([1, 28, 28])
```

```
1 x[0,6:14,6:14]
```

```
tensor([[0.0000, 0.0000, 0.1176, 0.1412, 0.3686, 0.6039, 0.6667, 0.9922],
        [0.0000, 0.1922, 0.9333, 0.9922, 0.9922, 0.9922, 0.9922, 0.9922],
        [0.0000, 0.0706, 0.8588, 0.9922, 0.9922, 0.9922, 0.9922, 0.9922],
        [0.0000, 0.0000, 0.3137, 0.6118, 0.4196, 0.9922, 0.9922, 0.8039],
        [0.0000, 0.0000, 0.0000, 0.0549, 0.0039, 0.6039, 0.9922, 0.3529],
        [0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.5451, 0.9922, 0.7451],
        [0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0431, 0.7451, 0.9922],
        [0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.1373, 0.9451]])
```

```
1 import matplotlib.pyplot as plt
2 plt.imshow(x[0], cmap="gray")
```

```
<matplotlib.image.AxesImage at 0x7f471c18c390>
```



Practice Problems 2, 3 & 4:

Minimizing a loss function using SGD

# Reporting test set loss

```
1 beta_hat = torch.ones((p, 1), requires_grad=True)
2 optimizer = torch.optim.SGD([beta_hat,], lr=2e-1)
3
4 for epoch in range(5):
5     train_mse = 0.0
6     for i, (x, y) in enumerate(train_dl):
7         mse = loss(x, y, beta_hat)
8         mse.backward()
9         optimizer.step()
10        beta_hat.grad.zero_()
11        train_mse += mse.item()
12
13    test_mse = 0.0
14    for i, (x, y) in enumerate(test_dl):
15        mse = loss(x, y, beta_hat)
16        test_mse += mse.item()
17
18    train_mse /= len(train_dl)
19    test_mse /= len(test_dl)
20
21    print("train_mse = %.3f    test_mse = %.3f" % (train_mse, test_mse))
```

```
train_mse = 0.248    test_mse = 0.000
train_mse = 0.000    test_mse = 0.000
train_mse = 0.000    test_mse = 0.000
train_mse = 0.000    test_mse = 0.000
train_mse = 0.000    test_mse = 0.000
```

# Refactoring

```
1 beta_hat = torch.ones((p, 1), requires_grad=True)
2 optimizer = torch.optim.SGD([beta_hat,], lr=2e-1)
3
4 def train_loop():
5     train_mse = 0.0
6     for i, (x, y) in enumerate(train_dl):
7         mse = loss(x, y, beta_hat)
8         mse.backward()
9         optimizer.step()
10        beta_hat.grad.zero_()
11        train_mse += mse.item()
12
13    return train_mse / len(train_dl)
14
15 def test_loop():
16     test_mse = 0.0
17     for i, (x, y) in enumerate(test_dl):
18         mse = loss(x, y, beta_hat)
19         test_mse += mse.item()
20
21    return test_mse / len(test_dl)
22
23 for epoch in range(5):
24     train_mse = train_loop()
25     test_mse = test_loop()
26     print("train_mse = %.3f    test_mse = %.3f" % (train_mse, test_mse))
```

# Modules

```
1 from torch import nn
2
3
4 class LinearRegression(nn.Module):
5
6     def __init__(self, p):
7         super(LinearRegression, self).__init__()
8         self.beta_hat = nn.Parameter(torch.ones((p, 1)))
9
10    def forward(self, x):
11        yhat = torch.mm(x, self.beta_hat)
12        return yhat
13
14
15 lr = LinearRegression(p)
16 optimizer = torch.optim.SGD(lr.parameters(), lr=2e-1)
17 getmse = torch.nn.MSELoss()
18
19 def train_loop():
20     train_mse = 0.0
21     for i, (x, y) in enumerate(train_dl):
22         yhat = lr(x)
23         mse = getmse(y, yhat)
24         lr.zero_grad()
25         mse.backward()
26         optimizer.step()
27         beta_hat.grad.zero_()
28         train_mse += mse.item()
29
30     return train_mse / len(train_dl)
```



# Built-in Linear modules

These two blocks are equivalent:

```
3 class LinearRegression(nn.Module):
4
5     def __init__(self, p):
6         super(LinearRegression, self).__init__()
7         self.beta_hat = nn.Parameter(torch.ones((p, 1)))
8
9     def forward(self, x):
10        yhat = torch.mm(x, self.beta_hat)
11        return yhat
12
13
14 lr = LinearRegression(p)
```

```
3 lr = nn.Linear(p, 1, bias=False)
```

# Nonlinear data fit by a linear model

```
1 n = 10000
2 p = 5
3 dataset_x = torch.randn(n, p)
4 dataset_y = (dataset_x ** 4).sum(1)
```

```
19 model = Linear(p, 1)
20 criterion = MSELoss()
21 optimizer = torch.optim.SGD(model.parameters(), lr=1e-2)
```

# Nonlinear data fit by a linear model


```
53 for epoch in range(1000):  
54     train_mse = train_loop()  
55     test_mse = test_loop()  
56     print("train_mse = %.3f  test_mse=%.3f" % (train_mse, test_mse))  
57
```

```
train_mse = 561.076  test_mse=496.446  
train_mse = 488.869  test_mse=484.272  
train_mse = 484.442  test_mse=483.528  
train_mse = 484.006  test_mse=483.461  
train_mse = 484.199  test_mse=483.510  
train_mse = 484.396  test_mse=483.412  
train_mse = 484.245  test_mse=483.409  
train_mse = 484.343  test_mse=483.417  
train_mse = 484.450  test_mse=483.555  
train_mse = 484.334  test_mse=483.538  
train_mse = 484.409  test_mse=483.567  
train_mse = 484.254  test_mse=483.723  
train_mse = 484.212  test_mse=483.500
```

# Nonlinear data fit by a nonlinear model

```
14 hidden_dim = 32
15 model = Sequential(
16     Linear(p, hidden_dim),
17     ReLU(),
18     Linear(hidden_dim, 1),
19 )
```

ReLU  
(Rectified Linear Unit)



$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

# Nonlinear data fit by a nonlinear model

```
51 for epoch in range(1000):  
52     train_mse = train_loop()  
53     test_mse = test_loop()  
54     print("train_mse = %.3f  test_mse=%.3f" % (train_mse, test_mse))  
55
```

```
train_mse = 344.467  test_mse=237.223  
train_mse = 202.776  test_mse=159.583  
train_mse = 161.530  test_mse=140.376  
train_mse = 147.999  test_mse=160.638  
train_mse = 139.355  test_mse=128.781  
train_mse = 134.093  test_mse=116.675  
train_mse = 127.323  test_mse=129.027  
train_mse = 119.114  test_mse=106.170  
train_mse = 112.629  test_mse=100.606  
train_mse = 103.340  test_mse=90.888  
train_mse = 95.971   test_mse=90.992
```

...

```
train_mse = 3.553  test_mse=3.333  
train_mse = 3.392  test_mse=3.245
```